

Building a LLM-powered Equity Research Tool

A Mini Project report submitted
in the partial fulfilment of the requirements for the award of the degree of

Bachelor of Technology In Computer Science & Engineering (Artificial Intelligence and Machine Learning)

By

21071A6614	Challa Kalyan Ram
21071A6634	Medida Vijaya Sai
21071A6651	Prachet Yerramalla
22075A6601	Kolloju Chandan

Under the Guidance of

Dr. Durga Prasad Kavadi

Sr. Assistant Professor

CSE – AIML & IoT, VNR VJIET

Submitted to



DEPARTMENT OF

**CSE- (ARTIFICIAL INTELLIGENCE AND MACHINE LEARNING &
INTERNET OF THINGS)**

**Vallurupalli Nageswara Rao Vignana Jyothi Institute of Engineering and Technology,
Hyderabad, Telangana**

August 2024



VALLURUPALLI NAGESWARA RAO VIGNANA JYOTHI INSTITUTE OF ENGINEERING AND TECHNOLOGY

An Autonomous, ISO 9001:2015 & QS I-Gauge Diamond Rated Institute, Accredited by NAAC with 'A++' Grade

NBA Accreditation for B.Tech. CE, EEE, ME, ECE, CSE, EIE, IT Programmes

Approved by AICTE, New Delhi, Affiliated to JNTUH, NIRF 135th Rank in Engineering Category

Recognized as "College with Potential for Excellence" by UGC

VignanaJyothi Nagar, Pragathi Nagar, Nizampet (S.O), Hyderabad – 500 090, TS, India.

Telephone No: 040-2304 2758/59/60, Fax: 040-23042761

CERTIFICATE

This is to certify that **CHALLA KALYAN RAM (21071A6614), MEDIDA VIJAYA SAI (21071A6634), PRACHET YERRAMALLA (21071A6651), KOLLOJU CHANDAN (22075A6601)** have successfully completed their Mini project work at CSE-(AIML & IoT) Department of VNRVJIET, Hyderabad entitled **“Building a LLM-powered Equity Research Tool”** in partial fulfilment of the requirements for the award of B. Tech degree during the academic year 2023-2024.

This work is carried out under my supervision and has not been submitted to any other University/Institute for award of any degree/diploma.

GUIDE

Dr. Durga Prasad Kavadi
Sr. Assistant Professor
Dept. of CSE (AIML & IoT)

Head of the Department

Dr. Sagar Yeruva
Associate Professor
Dept. of CSE (AIML & IoT)

DECLARATION

This is to certify that our project titled “**BUILDING A LLM-POWERED EQUITY RESEARCH TOOL**” submitted to Vallurupalli Nageswara Rao Vignana Jyothi Institute of Engineering and Technology in complete fulfilment of the requirement for the award of Bachelor of Technology in CSE- (Artificial Intelligence and Machine Learning) is a bonafide report to the work carried out by us under the guidance and supervision of Dr. Durga Prasad Kavadi, Sr. Assistant Professor, Department of CSE-(AIML & IoT), Vallurupalli Nageswara Rao Vignana Jyothi Institute of Engineering and Technology. To the best of our knowledge, this has not been submitted in any form to another University/Institute for an award of any degree/diploma.

Challa Kalyan Ram
21071A6614
Dept. of CSE (AIML&IoT)

Medida Vijaya Sai
21071A6634
Dept. of CSE (AIML&IoT)

Prachet Yerramalla
21071A6651
Dept. of CSE (AIML&IoT)

Kolloju Chandan
22075A6601
Dept. of CSE (AIML&IoT)

ABSTRACT

This project proposes an innovative tool to streamline equity research analysis by harnessing cutting-edge technologies. Equity analysts often grapple with the daunting task of sifting through numerous articles to gather relevant information. To tackle this challenge, our solution integrates Hugging Face models, Langchain, and Streamlit frameworks to develop a user-friendly news research tool. This tool empowers analysts to upload articles and pose inquiries, thereby saving time and enhancing accuracy in research endeavors. Technical intricacies involve leveraging Hugging Face models for question answering, Langchain for article processing, and Streamlit for UI development. Future iterations will include automated article scraping and integration with a vector database for heightened efficiency. By amalgamating advanced technologies, our project aims to revolutionize equity research analysis, offering substantial efficiency gains and improved research outcomes.

TABLE OF CONTENTS

I.	Table of Contents	1
II.	Table of Figures	3
III.	List of Tables	4
1	INTRODUCTION	5
1.1	EXISTING SYSTEM	5
1.2	PROPOSED SYSTEM	5
2	LITERATURE SURVEY	7
2.1	RELATED WORK	7
2.2	SYSTEM STUDY	9
3	DESIGN	10
3.1	SYSTEM REQUIREMENTS	10
3.1.1	Software Requirements	10
3.1.2	Hardware Requirements	11
3.2	UML	12
3.2.1	Use case	12
3.2.2	Activity Diagram	13
3.2.3	Class Diagram	14
4	IMPLEMENTATION	15
4.1	MODULES	15
4.1.1	Dataset	15
4.1.2	Models	15
4.1.3	Performance Metrics	16
4.2	OVERVIEW TECHNOLOGY	16
4.2.1	Data Collection	16
4.2.2	Data Processing	16
4.2.3	Algorithm selection	17
		1

4.2.4	User Interface	17
5	TESTING	18
5.1	TEST CASES	18
5.2	TEST RESULTS	19
6	RESULTS	20
6.1	SEMANTIC SEARCH AND EMBEDDING PERFORMANCE	20
6.1.1	Model Performance	20
6.1.2	Semantic Search Efficiency	21
6.1.3	Contextual Analysis	21
7	CONCLUSION	23
8	FUTURE SCOPE	24
8.1	INTEGRATION WITH REAL TIME DATA SOURCES	24
8.2	ADVANCED NATURAL LANGUAGE UNDERSTANDING	24
8.3	EXPANTION TO MULTILINGUAL SUPPORT	24
8.4	ENHANCED UI AND UX	24
9	BIBILOGRAPHY	25

TABLE OF FIGURES

Figure 1 Block Diagram	6
Figure 2 Use Case Diagram	12
Figure 3 Activity Diagram	13
Figure 4 Class Diagram	14
Figure 5 Time Vs Efficiency	22
Figure 6-1 Query 1	22
Figure 6-2 Query 2	23

LIST OF TABLES

Table 1 Test Cases	18
Table 2 Comparision of Processing Time with and with out caching	21

1 INTRODUCTION

Our project aims to revolutionize information extraction and comprehension from user-provided article links through web scraping. By parsing articles into digestible chunks and converting them into sentence embeddings, we facilitate efficient storage and retrieval of data. Utilizing semantic search techniques, users can pose questions related to the content, allowing for precise retrieval of relevant information. Integrated with pre-trained question-answering models, our system provides accurate answers and confidence scores, enhancing understanding and knowledge extraction. This project represents a fusion of web scraping, NLP, and question-answering technologies, empowering users with streamlined access to valuable insights within articles.

1.1 EXISTING SYSTEM:

Traditional Research:

- ❖ **Manual Research:** This involves manually sifting through financial reports, news articles, and other sources to find relevant information. This is time-consuming and prone to human error.
- ❖ **Investment Banking Research:** Large investment banks employ dedicated equity research analysts who provide research reports and recommendations to institutional investors. However, this service can be expensive and is not readily available to all investors.
- ❖ **Financial Databases:** Services like Bloomberg and S&P Capital IQ offer comprehensive financial data and research reports on publicly traded companies. These can be expensive subscriptions and require some level of financial literacy to navigate effectively.

1.2 PROPOSED SYSTEM:

Our project proposes a compelling solution to address the challenges faced by equity analysts in their research process. Let's delve into the key aspects of our proposed system:

1.2.1 Core Functionalities:

- **Information Processing:** Utilizes Langchain to automate article processing, potentially extracting key information like entities, relationships, and sentiment.
- **Question Answering:** Leverages BERT to enable analysts to pose specific questions about the processed information from articles.
- **User-Friendly Interface:** Employs Streamlit to create a user-friendly platform for uploading articles and interacting with the question-answering capabilities.

1.2.2 Potential Benefits:

- **Improved Efficiency:** By automating information gathering and answering specific questions, analysts can save significant time and effort.
- **Enhanced Accuracy:** Extracting key information and answering specific queries can lead to a more precise understanding of the information contained in the articles, potentially improving research accuracy.
- **Accessibility:** The user-friendly interface can potentially make the tool accessible to a wider range of users, including individual investors or smaller firms.

1.2.3 Future Enhancements:

- **Automated Article Scraping:** Integrating scraping functionalities can further reduce manual work by automatically collecting relevant articles for analysis.
- **Vector Database Integration:** Implementing a vector database might enable efficient storage and retrieval of processed information, improving search and retrieval capabilities.




Figure 1 Block Diagram

2 LITERATURE SURVEY

2.1 RELATED WORK

We referenced several research papers focusing on various aspects of document summarization and language models. These studies address different challenges, methodologies, and Advancements in the field.

2.1.1 FlexiGPT: Engaging with Documents (2024)

This paper addresses the inefficiency of existing embedding models in extracting top results from LLM knowledge bases based on user prompts. The objective was to leverage LLMs to pre-train text summarization models, enhancing the extraction of relevant results from documents. The strengths of this approach include its ability to store and tokenize data effectively for accurate answer retrieval. However, the model's reliance on tokenized data may limit its effectiveness in handling diverse and complex queries.

2.1.2 An Empirical Survey on Long Document Summarization: Datasets, Models, and Metrics (2023)

The research provides a systematic evaluation of long document summarization, focusing on benchmark datasets, summarization models, and evaluation metrics. The study's strengths lie in its comprehensive review of model architectures and summarization strategies. However, the paper's focus on long documents might overlook advancements in summarizing shorter texts or different types of documents.

2.1.3 Text Document Summarization Using Word Embedding (2022)

This survey introduces Multi-document Summarization (MDS) using neural networks, addressing the challenge of summarizing multiple documents simultaneously. The proposed taxonomy and network design strategies highlight the advancements in MDS. Nevertheless, the survey may not fully address the limitations of current summarization models and their practical applications.

2.1.4 Abstractive Long Text Summarization Using Large Language Models (2023)

The paper presents a novel method for summarizing complex PDFs using LLMs, focusing on breaking down data into clusters and summarizing them. The strengths include the reduction of redundant data and improved summarization accuracy. However, challenges remain in handling very large or highly complex documents effectively.

2.1.5 Towards a Robust Retrieval-Based Summarization System (2024)

This research introduces SummRAG, a system designed to enhance the robustness of LLMs in summarization tasks. The approach focuses on evaluating and improving the robustness of LLMs through structured testing methods. While effective, the study's applicability might be limited by the specific scenarios used for evaluation.

2.1.6 Text Document Summarization Using Word Embedding (2019)

This paper proposes a summarization technique that emphasizes semantic preservation using distributional semantic models. The strengths include improved precision and scalability. However, the computational cost of semantic models and their impact on summary quality are noted limitations.

2.1.7 MeaningBERT: Assessing Meaning Preservation Between Sentences (2023)

MeaningBERT introduces a metric for assessing meaning preservation between sentences using BERT-Base. The strengths include its ability to correlate well with human judgment. However, the metric's reliance on specific training data and its potential limitations in complex tasks are drawbacks.

2.1.8 Natural Language Question Answering System as a Participant in Human Q&A Portals (2011)

This paper focuses on the Wrong Answer Avoidance (WAA) method to prevent automated systems from providing incorrect answers. The method's strengths lie in its machine learning approach to question classification. However, it may face challenges related to the diversity of questions and the accuracy of predictions.

2.1.9 Semantic Search on Text and Knowledge Bases (2023)

The paper provides an overview of semantic search techniques, offering a foundational understanding of concepts, techniques, and future directions. Its strengths include a comprehensive review of semantic search advancements. However, it might not cover the latest innovations or specific applications in depth.

2.1.10 Platypus: Quick, Cheap, and Powerful Refinement of LLMs (2023)

Platypus introduces a cost-effective method for LLM fine-tuning using less data and computational resources. The strengths include its efficiency and strong performance on benchmarks. Limitations include the potential challenges in evaluating the largest LLM models due to computational constraints.

2.1.11 Learning from Mistakes Makes LLM Better Reasoner (2023)

The paper proposes LEMA, a novel approach incorporating mistake-correction data to improve LLM reasoning capabilities. The strengths include enhanced performance in reasoning tasks. However, the method's reliance on specific mistake-correction data may limit its generalizability.

2.1.12 LLM-Adapters: An Adapter Family for Parameter-Efficient FineTuning of Large Language Models (2023)

This paper explores parameter-efficient fine-tuning (PEFT) methods for LLMs, focusing on LLM-Adapters for simplifying PEFT implementation. The strengths include effective fine-tuning with smaller models. Limitations include the inability to test the largest LLM models due to computational constraints.

2.1.13 Blended RAG: Improving RAG Accuracy with Semantic Search and Hybrid Query-Based Retrievers (2024)

Blended RAG combines semantic search with hybrid query strategies to improve RAG system accuracy. The strengths include enhanced retrieval performance and adaptability. However, challenges such as missing benchmarks and evaluation metrics remain.

2.1.14 Vector Database Management Systems: Fundamental Concepts, Use-Cases, and Current Challenges (2024)

This paper provides an introduction to vector database management systems (VDBMS), covering their functionalities, use cases, and challenges. The strengths include a comprehensive overview of VDBMS applications. Limitations involve challenges in balancing speed and accuracy.

2.2 SYSTEM STUDY

In this comprehensive study, various papers were analysed to understand advancements in text summarization, LLMs, and semantic search. The primary focus was on evaluating and improving methodologies related to document summarization, model efficiency, and cost-effectiveness in processing large-scale data. Each paper provided insights into specific challenges and innovations, highlighting the ongoing progress and areas for future research in these fields.

3 DESIGN

3.1 SYSTEM REQUIREMENTS

3.1.1 Software Requirements:

3.1.1.1 Programming Languages:

- **Python:** Essential for implementing the script, handling data processing, and managing model inference.

3.1.1.2 Web Scraping and Document Loading:

- **langchain_community:** Required for loading unstructured data from URLs.
- **UnstructuredURLLoader:** Utilized for fetching and loading data from the provided URLs.

3.1.1.3 Text Splitting and Preprocessing:

- **langchain.text_splitter (RecursiveCharacterTextSplitter):** Used for splitting text into manageable chunks, essential for efficient processing and embedding.

3.1.1.4 NLP and Machine Learning Libraries:

- **Transformers:** Necessary for using the pipeline function to create a question-answering model.
- **Sentence-Transformers:** Required for embedding the text chunks and the query using the SentenceTransformer model.

3.1.1.5 Semantic Search Tools:

- **sentence-transformers/util:** Utilized for performing semantic search, matching the query with the most relevant text chunk.

3.1.1.6 Data Storage and Caching:

- **pickle:** For serializing and deserializing Python object structures to cache embeddings and text chunks.
- **hashlib:** Utilized for generating a unique cache key based on URLs.

3.1.1.7 Data Storage and Management:

- **Local Storage:** To store and retrieve cached data, including document chunks and embeddings.

3.1.2 Hardware Requirements:

3.1.2.1 Computing Power:

- **Multi-Core CPU:** Necessary for handling the text processing, embedding generation, and semantic search operations efficiently

3.1.2.2 GPU:

- Having access to a GPU (Graphics Processing Unit) can significantly speed up training deep learning models. Recommended for accelerating model inference tasks, especially when using the transformers pipeline for question-answering.

3.1.2.3 RAM:

- Adequate RAM, typically 16 GB or more, is essential for handling large datasets and training complex models.

3.1.2.4 Storage:

- Sufficient storage space for datasets and model checkpoints. SSDs are preferable for faster data access.

3.1.2.5 Internet Connection:

- A reliable internet connection is necessary for downloading pre-trained models (like the ones from Hugging Face), as well as for accessing and loading data from the provided URLs.

3.2 UML

3.2.1 Use case:

They are usually used to illustrate the various actions taken by the application. They also show the several users who can carry out these functions. Use-case diagrams fall under behaviour diagrams due to their emphasis on the tasks carried out and the users (actors) who carry out these tasks.




Figure 2 Use Case Diagram

3.2.2 Activity Diagram:

An activity diagram visually presents a series of actions or flow of control in a system like a flowchart or a data flow diagram. Activity diagrams are often used in business process modelling.




Figure 3 Activity Diagram

3.2.3 Class Diagram

Class diagram is a static diagram. It represents the static view of an application. Class diagram is not only used for visualizing, describing, and documenting different aspects of a system but also for constructing executable code of the software application.




Figure 4 :Class Diagram

4 IMPLEMENTATION

4.1 MODULES

4.1.1 Dataset:

The dataset is the foundational element upon which the entire system is built. It comprises a collection of articles obtained through web scraping, which are then processed to extract meaningful content. These articles are sourced from user-provided links and are loaded into the system using the *UnstructuredURLLoader* module from the *langchain_community* package. Given the unstructured nature of the data, the text is divided into manageable chunks using the *RecursiveCharacterTextSplitter*, facilitating subsequent processing and embedding. This segmentation is essential for efficiently handling and analysing the large volumes of textual data, ensuring that the system can operate smoothly without being overwhelmed by the sheer quantity of information.

4.1.2 Models

The models developed for this project are central to its ability to process, analyse, and retrieve information. Two primary types of models were implemented: Machine Learning (ML) models and Deep Learning (DL) models, each contributing different capabilities to the system.

4.1.2.1 Machine Learning Models

4.1.2.1.1 Sentence-Transformers

The ML models focus on the initial stages of data processing, specifically the transformation of text chunks into sentence embeddings. The model used for this purpose is the Sentence-Transformers model, which is designed to convert text into high-dimensional vectors that capture the semantic meaning of the content. These embeddings are crucial for enabling the system's semantic search capabilities, allowing it to map user queries to the most relevant chunks of text. By leveraging these embeddings, the system can effectively filter and prioritize information, ensuring that users receive the most pertinent responses to their queries.

4.1.2.2 Deep Learning Network

4.1.2.2.1 Pre-trained Transformers (BERT, QA Model)

The DL models are employed in the later stages of the process, where the system must generate precise answers to user queries. For this, the system utilizes pre-trained transformer models, such as BERT, which are known for their advanced natural language processing capabilities. These models are integrated into the question-answering pipeline, where they analyse the text chunks identified by the semantic search process and generate accurate and contextually relevant responses. The deep learning models are particularly effective in understanding the nuances of language, allowing the system to provide high-quality answers even in complex or ambiguous scenarios.

4.1.3 Performance Metrics

Performance metrics provide a quantitative assessment of model effectiveness. Key metrics include accuracy, measuring overall correctness; precision, indicating true positive rate; recall, gauging the ability to identify positives; and F1 score, blending precision and recall.

- **Accuracy:** Measures the ratio of correctly predicted instances to the total instances.
Formula: $(TP + TN) / (TP + TN + FP + FN)$
- **Precision:** Reflects the proportion of true positives among the predicted positive instances.
Formula: $TP / (TP + FP)$
- **Recall (Sensitivity or True Positive Rate):** Represents the ratio of true positives among actual positive instances.
Formula: $TP / (TP + FN)$
- **F1 Score:** Balances precision and recall, offering a harmonic mean of the two metrics. Formula: $2 * (Precision * Recall) / (Precision + Recall)$

4.2 OVERVIEW TECHNOLOGY

The technology stack chosen for this project is integral to its success, enabling the system to handle complex tasks such as web scraping, text processing, semantic search, and question answering. This section outlines the specific tools and frameworks used at each stage of the implementation, highlighting how they contribute to the overall functionality of the system.

4.2.1 Data Collection

Data collection is the first step in the system's workflow, involving the extraction of articles from the web based on user-provided URLs. The process begins with the utilization of the *UnstructuredURLLoader*, a tool from the *langchain_community* package, which is designed to efficiently load unstructured data from web pages. This tool enables the system to access a wide range of online content, ensuring that users can retrieve information from diverse sources. Once the data is loaded, it is parsed and stored locally, ready for further processing.

4.2.2 Data Processing

The data processing stage is where the raw textual content is transformed into a structured format that can be easily analysed and queried. Initially, the text is divided into smaller, more manageable chunks using the *RecursiveCharacterTextSplitter*, which ensures that each chunk is of optimal size for embedding and further analysis. These chunks are then converted into sentence embeddings using the Sentence-Transformers model. This process involves mapping the textual data into a high-dimensional space where semantically similar pieces of text are positioned close to one another. This transformation is crucial for enabling the system's semantic search capabilities, which allow it to match user queries with the most relevant chunks of information.

4.2.3 Algorithm selection – Question-Answering

The question-answering module is the core of the system, where user queries are processed and matched with relevant information from the dataset. The process begins with the user inputting a query, which is then converted into an embedding using the same Sentence-Transformers model employed during data processing. The system performs a semantic search to identify the text chunks most closely aligned with the query. These chunks are then passed through a pre-trained transformer model, such as QA Model, which analyses the content and generates a precise answer to the query. This step ensures that the system can deliver accurate and contextually appropriate responses, enhancing the user's ability to extract meaningful insights from the data.

4.2.4 User Interface

The user interface (UI) is designed to provide an intuitive and user-friendly experience, allowing users to interact with the system without requiring extensive technical knowledge. The UI is built using Streamlit, a framework known for its simplicity and effectiveness in creating interactive web applications. Through the Streamlit interface, users can easily upload article links, input queries, and view the system's responses. The interface also displays confidence scores alongside the answers, giving users additional context to assess the reliability of the information provided. This focus on usability ensures that the system is accessible to a wide range of users, including those who may not have a deep background in data science or machine learning.

5 TESTING

5.1 TEST CASES

Table 1 Test Cases

Test Case ID	Test Case Description	Test Steps	Preconditions	Test Data	Post Conditions
TC_QA_01	Answer question about Ola Electric	1. Import URLs 2. Load and split text 3. Encode text 4. Perform semantic search 5. Generate and return answer	URLs must be accessible, models loaded	Question: "What is Ola Electric?" URLs: 1. URL 1 2. URL 2	Relevant context found and correct answer returned with confidence score
TC_QA_02	Test with an unrelated question	1. Import URLs 2. Load and split text 3. Encode text 4. Perform semantic search 5. Generate and return answer	URLs must be accessible, models loaded	Question: "What is the capital of France?" URLs: 1. URL 1 2. URL 2	No relevant context found, appropriate message displayed
TC_Cache_01	Validate cache functionality	1. Run project with initial URLs 2. Re-run project with the same URLs	Cache should be enabled	URLs: 1. URL 1 2. URL 2	Chunks and embeddings loaded from cache on the second run
TC_Cache_02	Validate cache invalidation	1. Run project with initial URLs 2. Modify URL list and re-run project	Cache should be enabled	URLs: 1. URL 1 2. URL 2	New cache created, new chunks and embeddings processed
TC_Encode_01	Test sentence embedding	1. Load data 2. Split text 3. Encode text to embeddings	Model 'all-MiniLM-L6-v2' must be loaded	Text: "Ola Electric IPO was subscribed 1.48 times on Day 1 of the issue."	Embeddings generated successfully

5.2 TEST RESULTS

The test cases were designed to evaluate various functionalities of the question-answering (QA) system, including its ability to handle specific questions, utilize caching effectively, and perform semantic searches. The tests also assessed the system's robustness in handling errors and its overall performance with different inputs.

- **Test Case ID: TC_QA_01**
Description: Answer question about Ola Electric
Result: The model successfully answered the question "What is Ola Electric?" using the provided URLs. Relevant context was found and the answer was returned with a high confidence score.
- **Test Case ID: TC_QA_02**
Description: Test with an unrelated question
Result: The model handled the question "What is the capital of France?" appropriately, noting that no relevant context was available. It displayed an appropriate message indicating that no relevant context was found.
- **Test Case ID: TC_Cache_01**
Description: Validate cache functionality
Result: The cache functionality was validated by running the project with the same URLs. On the second run, the model successfully loaded chunks and embeddings from the cache, demonstrating effective cache usage.
- **Test Case ID: TC_Cache_02**
Description: Validate cache invalidation
Result: After modifying the URL list and re-running the project, the model created a new cache with updated chunks and embeddings, confirming that the cache invalidation and update mechanisms work correctly.
- **Test Case ID: TC_Encode_01**
Description: Test sentence embedding
Result: Sentence embeddings were successfully generated using the model 'all-MiniLM-L6-v2' for the text chunk, indicating proper functionality of the embedding process.

6 RESULTS

The project successfully implemented a question-answering system, showcasing its effectiveness in handling data and queries. Initially, data was imported from a variety of URLs, including sources on IPOs, market news, and SQL tutorials. The question “What is Ola Electric?” was also incorporated for testing.

A caching mechanism was utilized to enhance performance, with a unique cache key generated from the URLs. This allowed the system to quickly retrieve pre-processed data, minimizing redundant computations and speeding up response times.

The data was processed by splitting it into manageable chunks of 450 characters, effectively handling text segmentation with separators such as new lines and punctuation. This pre-processing step ensured that the text was organized into coherent segments suitable for analysis.

For semantic understanding, the text chunks were encoded using the SentenceTransformer model. The query was transformed into an embedding and used to perform a semantic search, identifying the most relevant chunk. Finally, the DistilBERT model was employed to generate answers, providing accurate and contextually relevant responses to the question. This comprehensive approach demonstrated the system's robustness in delivering precise answers based on the processed text.

6.1 SEMANTIC SEARCH AND EMBEDDING PERFORMANCE

6.1.1 Model Performance

Table 2: Comparison of processing time with and without caching.

Model	Processing Time (Without Cache)	Processing Time (With Cache)
Semantic Search Model	12.7 seconds	2.3 seconds

6.1.2 Semantic Search Efficiency

- **Caching Mechanism:** The implementation of the caching mechanism significantly reduced processing time from 12.7 seconds to 2.3 seconds, emphasizing the importance of caching in improving system efficiency.




Figure 5 Time Vs Efficiency

6.1.3 Contextual Analysis

- **Semantic Search Result:** The model identified the relevant context with a confidence score of 96.49%, showcasing its effectiveness in retrieving pertinent information from the corpus.
- **Question-Answering Accuracy:** The generated answer, "August 1, 2024" correctly matched the expected answer, demonstrating the model's precision.

News Article URLs

URL 1

<https://www.chittorgarh.com/ipo/ceigall>

URL 2

<https://hbswk.hbs.edu/item/after-silicon>

URL 3

<https://www.moneycontrol.com/news/b>

Equity Research Tool

Enter your question:

Ceigall India IPO bidding start date ?

Get Answer

Answer: August 1, 2024

Score: 96.49%

Figure 6-1 Query 1

- **Semantic Search Result:** The model identified the relevant context but with a confidence score of 6.24%, which indicates that the answer it provided was incorrect and highlights its ineffectiveness in retrieving pertinent information from the corpus
- **Question-Answering Accuracy:** The generated answer, "ceigall" correctly matched the expected answer, demonstrating the model's precision.

News Article URLs

URL 1

<https://www.chittorgarh.com/ipo/ceigall>

URL 2

<https://hbswk.hbs.edu/item/after-silicon>

URL 3

<https://www.moneycontrol.com/news/b>

Equity Research Tool

Enter your question:

Who is prime minister of india

Get Answer

Answer: ceigall

Score: 6.24%

Figure 6-2 Query 2

7 CONCLUSION

The development and implementation of the Equity Research Tool represent a significant step forward in the application of natural language processing (NLP) and machine learning techniques for automated question-answering tasks. Through the integration of various components such as URL data extraction, text chunking, semantic search, and contextual question answering, the project successfully demonstrates the ability to process unstructured web data and generate accurate answers to specific queries.

Key achievements of this project include:

1. **Data Handling and Caching:** The project efficiently handles multiple web-based data sources, splitting and storing text data in manageable chunks. The implementation of a caching mechanism ensures that repetitive processing is avoided, leading to faster response times for repeated queries.
2. **Semantic Search:** By leveraging the power of sentence transformers, the model performs semantic searches to identify the most relevant text chunks that match a given query. This approach enhances the model's ability to retrieve the most contextually appropriate information.
3. **Question-Answering Pipeline:** The use of a pre-trained transformer-based question-answering pipeline (BERT) allows for the extraction of precise answers from the identified text chunks, showcasing the potential of combining semantic search with sophisticated NLP models.

In conclusion, the Equity Research Tool demonstrates a practical approach to automated question answering by combining state-of-the-art NLP techniques with efficient data management strategies. The insights gained from this project lay the groundwork for further exploration and refinement in the domain of intelligent information retrieval and question answering systems.

8 FUTURE SCOPE

The future scope of this project is vast, with numerous opportunities to enhance functionality, scalability, and user experience. Below are some key areas where the project can be further developed:

8.1 INTEGRATION WITH REAL-TIME DATA SOURCES

Expanding the system to include real-time data scraping and processing could enhance its utility, especially for fields like finance and news analytics. This would allow users to retrieve the most current information, making the system more relevant for time-sensitive queries.

8.2 ADVANCED NATURAL LANGUAGE UNDERSTANDING (NLU)

Future iterations of the system could incorporate more sophisticated NLU models to better handle complex queries and context-dependent questions. This could improve the accuracy of answers in scenarios where deeper semantic understanding is required.

8.3 EXPANSION TO MULTILINGUAL SUPPORT

Adding multilingual capabilities would broaden the system's applicability across different languages, making it accessible to a global audience. This would involve integrating multilingual models and training the system on diverse datasets.

8.4 ENHANCED UI AND UX

Future developments could focus on improving the user interface, perhaps by incorporating voice-based queries, personalized recommendations, or interactive visualizations. These enhancements would make the system more intuitive and user-friendly, catering to a wider range of users.

9 BIBILOGRAPHY

1. “FlexiGPT: Engaging with Documents” by Abdalrhman Alquaary and Numan Çelebi (<https://doi.org/10.36287/setsci.6.1.029>)
2. “An Empirical Survey on Long Document Summarization: Datasets, Models and Metrics” by Huan Yee Koh, et al. (<https://doi.org/10.1145/3545176>)
3. “Text document summarization using word embedding” by Congbo Ma, et al. (<https://doi.org/10.1145/3529754>)
4. “Abstractive Long Text Summarization using Large Language Models” by Gunjan Keswani, et al. (<https://ijisae.org/index.php/IJISAE/article/view/4500>)
5. “Towards a Robust Retrieval-Based Summarization System” by [Shengjie Liu](#), et al. (<https://doi.org/10.48550/arXiv.2403.19889>)
6. “Text document summarization using word embedding” by Mudasir Mohd, et al. (<https://doi.org/10.1016/j.eswa.2019.112958>)
7. “MeaningBERT: assessing meaning preservation between sentences” by David Beauchemin, et al. ([10.3389/frai.2023.1223924](https://doi.org/10.3389/frai.2023.1223924))
8. “Natural Language Question Answering System as a Participant in Human Q&A Portals” by T Dong, et al. ([10.5591/978-1-57735-516-8/IJCAI11-405](https://doi.org/10.5591/978-1-57735-516-8/IJCAI11-405))
9. “Semantic Search on Text and Knowledge Bases” by Hannah Bast, et al. ([http://dx.doi.org/10.1561/15000000032](https://dx.doi.org/10.1561/15000000032))
10. “Platypus: Quick, Cheap, and Powerful Refinement of LLMs” by Ariel N. Lee, et al. (<https://doi.org/10.48550/arXiv.2308.07317>)
11. “Learning From Mistakes Makes LLM Better Reasoner” by Shengnan An, et al. (<https://doi.org/10.48550/arXiv.2310.20689>)
12. “LLM-Adapters: An Adapter Family for Parameter-Efficient Fine-Tuning of Large Language Models” by Zhiqiang Hu, et al. (<https://doi.org/10.48550/arXiv.2304.01933>)
13. “Blended RAG: Improving RAG (Retriever-Augmented Generation) Accuracy with Semantic Search and Hybrid Query-Based Retrievers” by Kunal Sawarkar, et al. (<https://doi.org/10.48550/arXiv.2404.07220>)
14. “Vector database management systems: Fundamental concepts, use-cases, and current challenges” by Toni Taipalus (<https://doi.org/10.1016/j.cogsys.2024.101216>)
15. “Cost-efficient domain-specific question answering using LLMs” by Md Adnan Arefeen, et al. (<https://doi.org/10.1016/j.nlp.2024.100065>)
16. “Natural Language Processing in finance : analysis of sentiment and complexity of news and earnings reports of swiss SMEs and their relevance for stock returns” by Schlaubitz and Alexander (<https://doi.org/10.21256/zhaw-24410>)